

National Research University Higher School of Economics
Faculty of Computer Science
Bachelor's Program "HSE University and University of London Double Degree
Program in Data Science and Business Analytics"

Introduction to Programming

Workshop #13

Wed 24.02.2021


Julio Carrasquel




NATIONAL RESEARCH
UNIVERSITY

Today's outline

`std:: [unordered_] [multi] maps and sets`



std::map and std::set- RBTrees



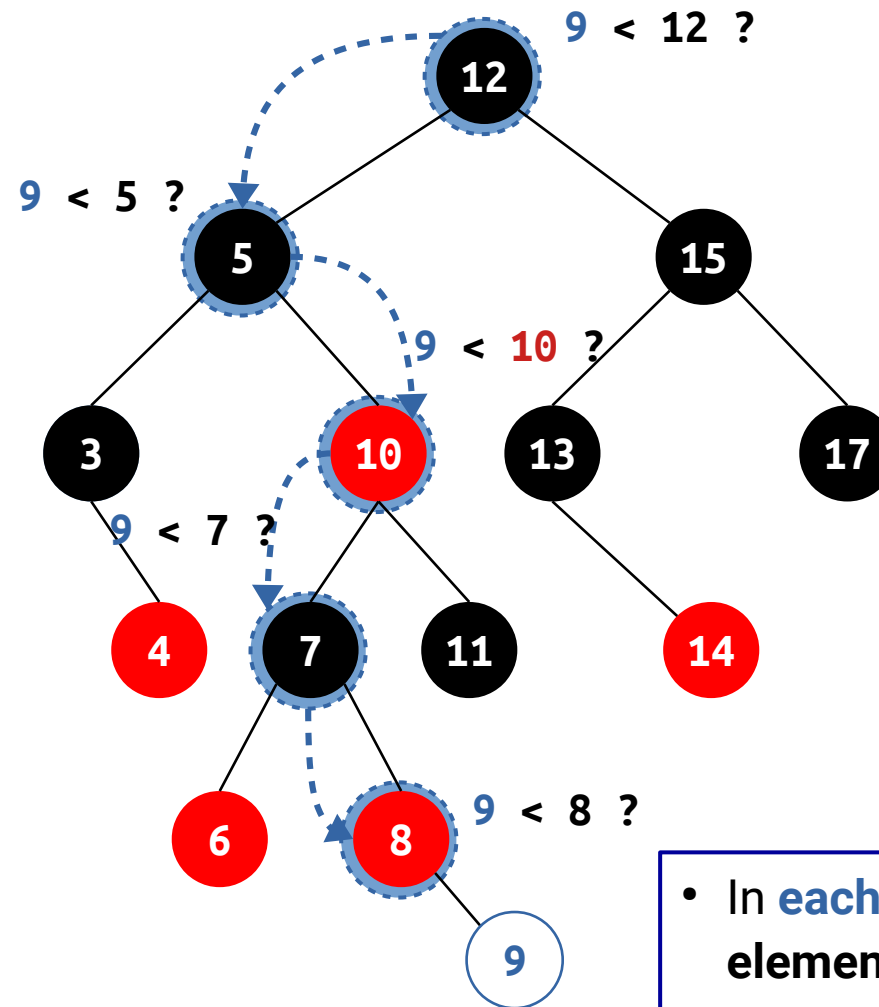
- `std::map` and `std::set` containers are *generally* implemented using Red-Black Trees (RBTrees)
- A RBTREE is a *self-balancing* binary search tree.
- A series of rules based on **red** and **black** colors allow to *maintain balance* in the tree.
- It has $O(\log n)$ for three basic operations: *find*, *insert* and *remove*.

Play with the example at:

<https://www.cs.usfca.edu/~galles/visualization/RedBlack.html>


RBTrees

Let's insert element 9



- In **each step**, we check if the **element to insert** *is less* than the **current element**.
- We need operator $<$ for that.

RBTrees with `std::string` as key




Print order: avocado berry cherry wiki watermelon

- By default, `std::less<std::string>` sort elements in the containers `std::map` and `std::set` in *lexicographical order*.
- But, we may create a custom comparator for strings.


RBTrees – **length** of a `std::string` as the key

Create a `std::set<std::string>` where elements are ordered based on the **length of a string**

We need to create a custom *compare function*.




avocado berry cherry kiwi watermelon



kiwi berry cherry avocado watermelon

RBTrees – **length** of a `std::string` as the key

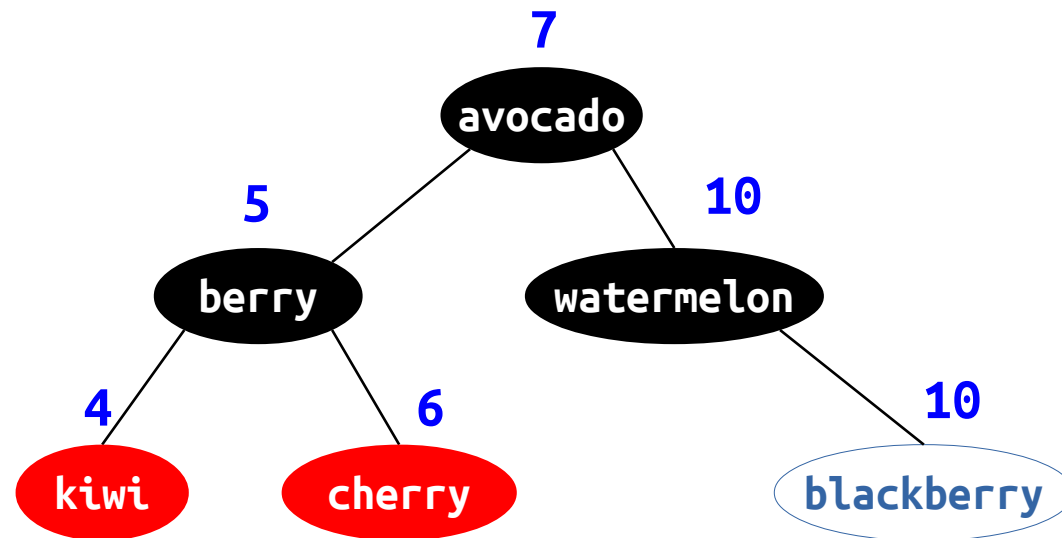
What if we want to add the string “blackberry” to the `std::set` organized by the string length?



kiwi berry cherry avocado watermelon

RBTrees – **length** of a `std::string` as the key

What if we want to add the string “blackberry” to the `std::set` organized by the string length? Use `std::multiset`



kiwi berry cherry avocado watermelon blackberry



We need to use `std::multiset` to allow multiple keys with the same value!

`std::multiset` (and `std::multimap`) are also implemented with RBTrees, as they still can handle multiple keys with the same value.

std::multimap - Example

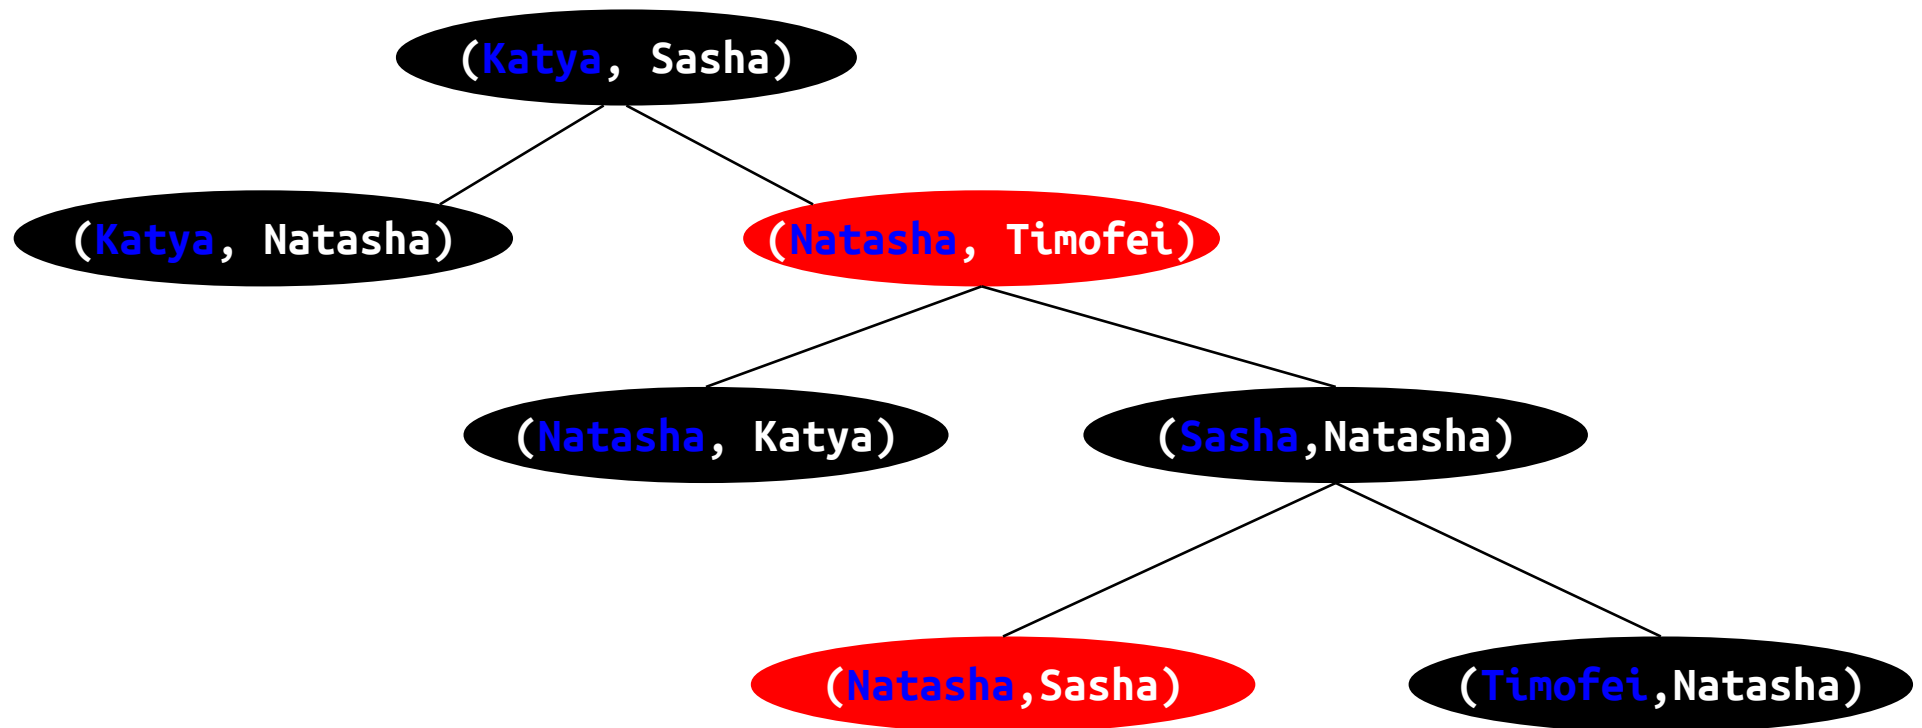
A `std::multimap` can be used to express relationships between elements.

"Who-follows-who" in a social network?




std::multimap - Example

Remember that `std::multimap` (and `std::map`) is also generally implemented as a RBTREE, where pairs (key, element) are ordered by the **key**.




std::multimap - Example

Remember that `std::multimap` (and `std::map`) is also generally implemented as a `RBTree`, where pairs (key, element) are ordered by the **key**.



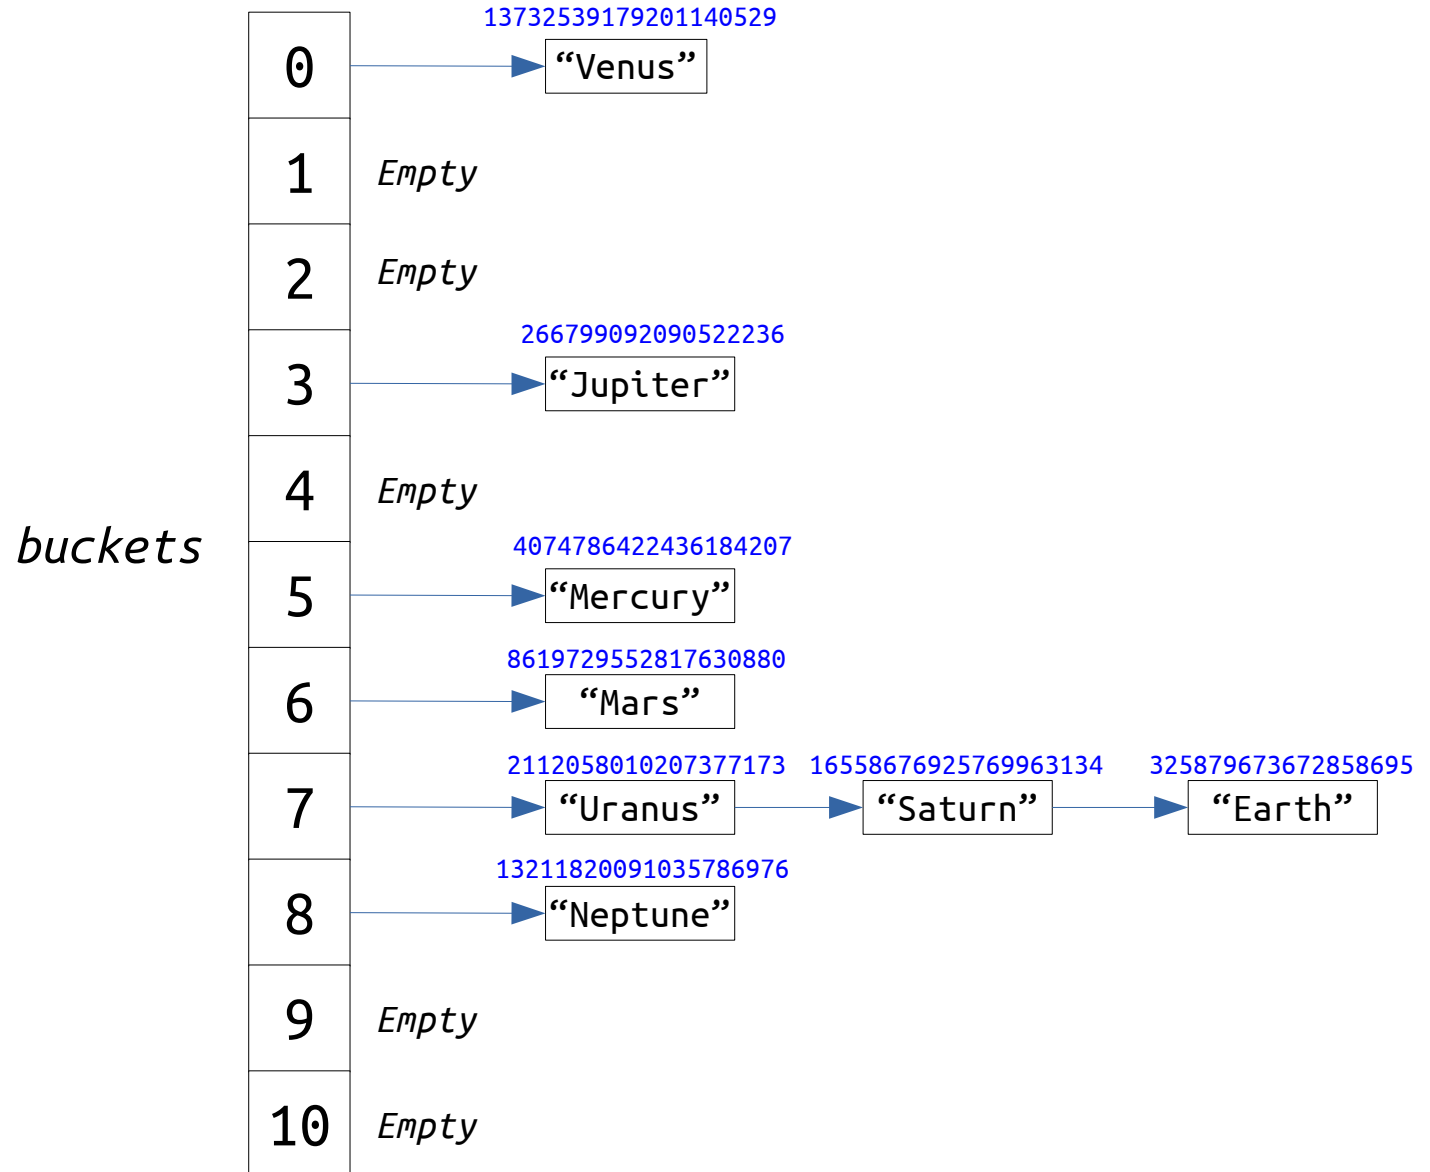
What if it is not important the internal ordering of elements in a `std::map`, `std::multimap`, `std::set`, or `std::multiset`? Let's use `std::unordered_` !

std:: [unordered_] [multi] maps and sets

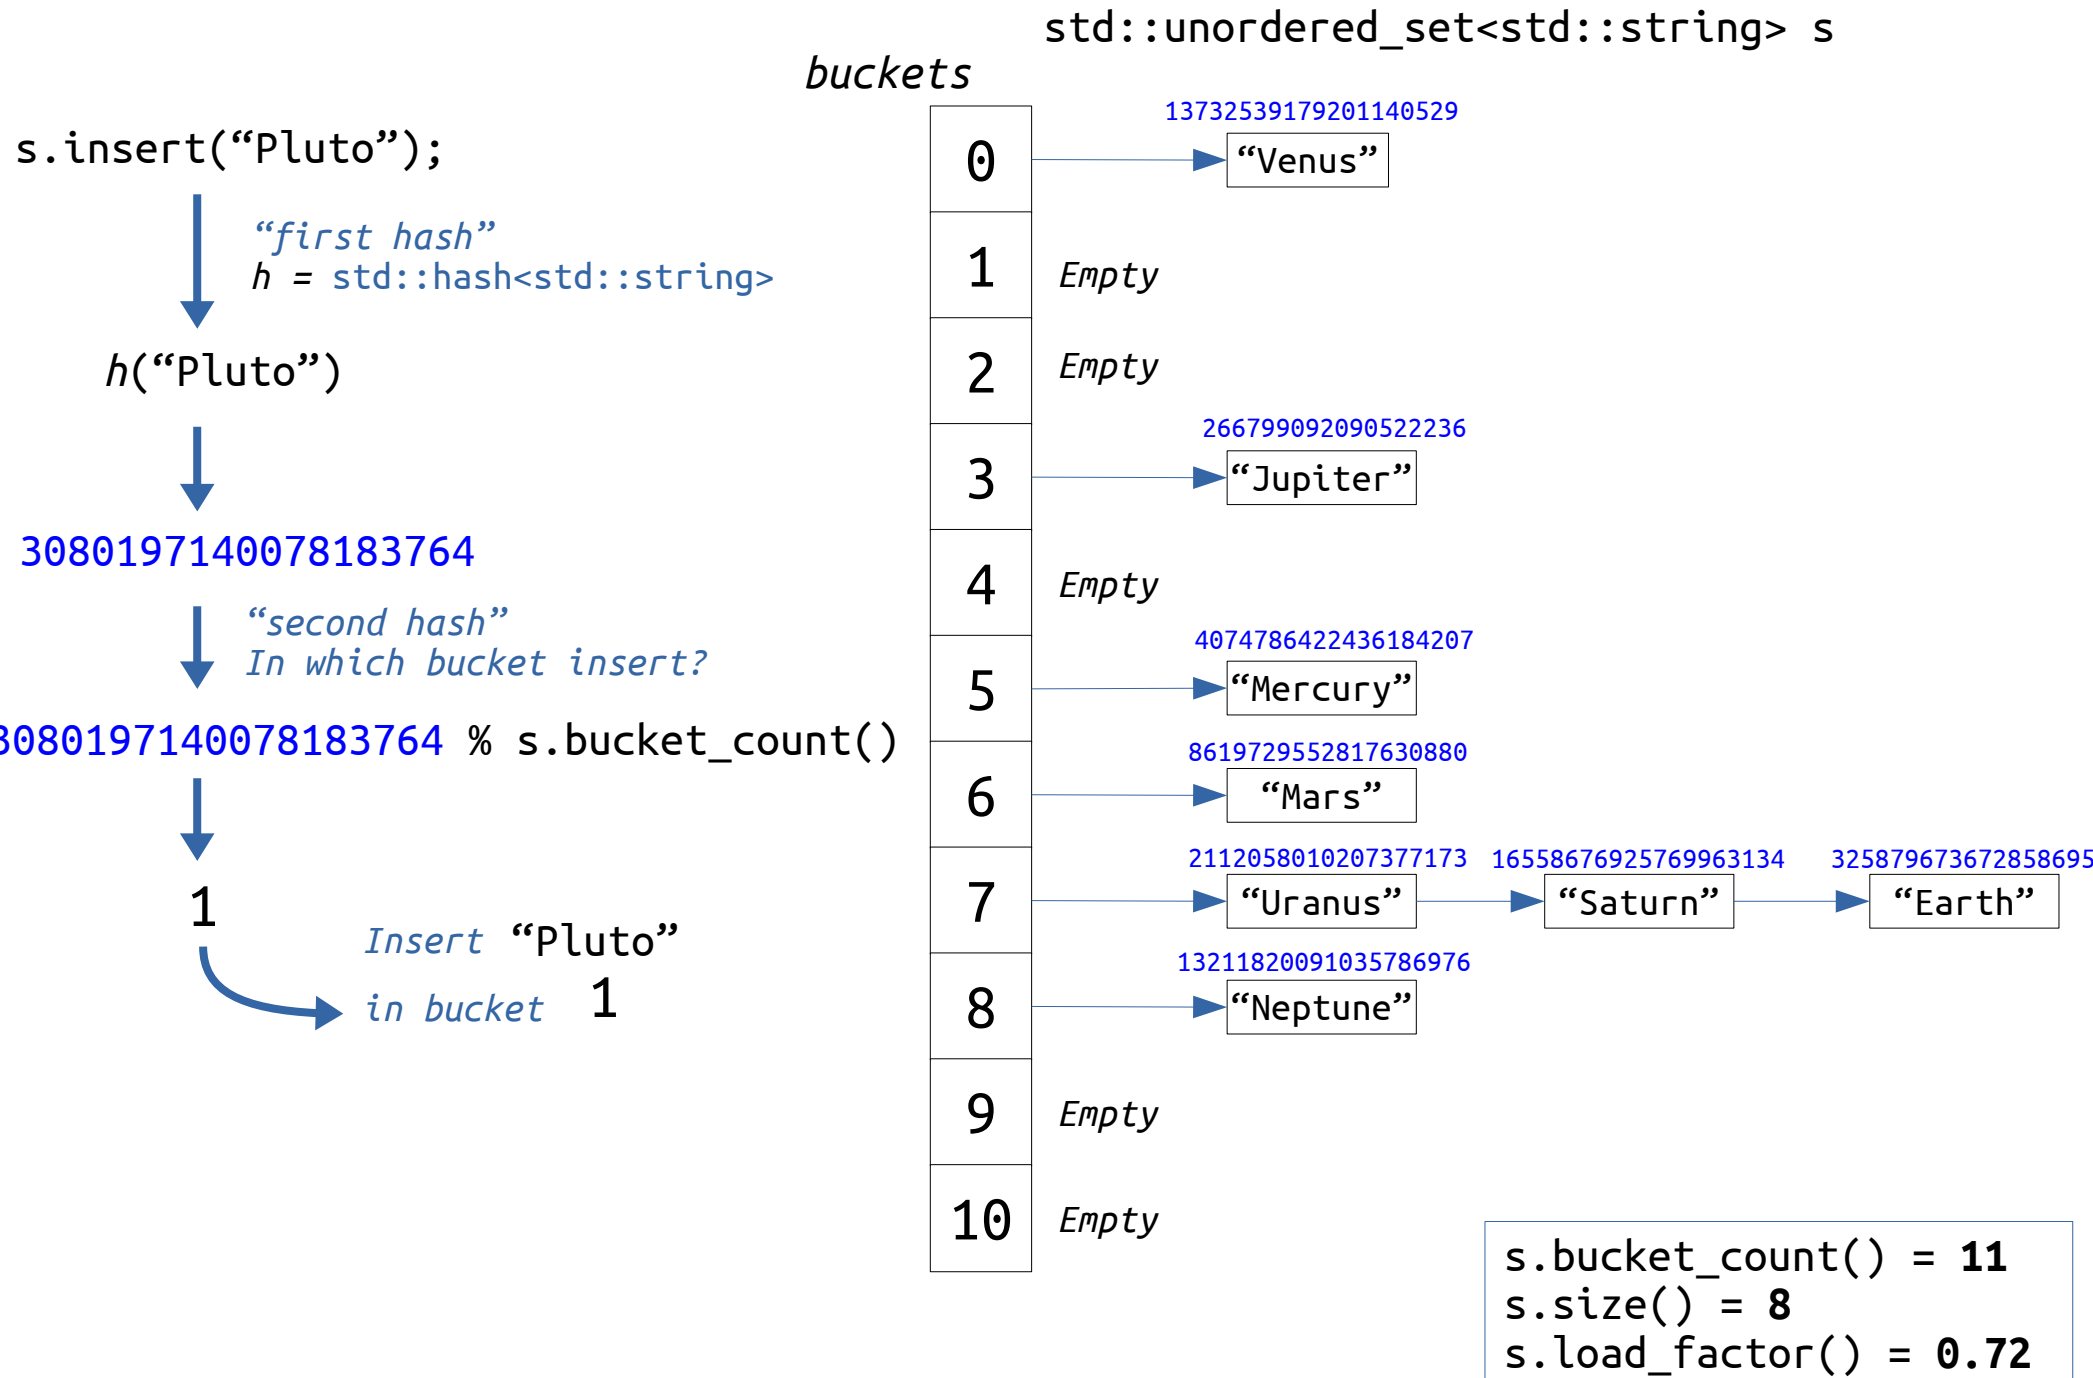


std::unordered_[multi] set/map and Hash functions

- Unordered containers are *generally* supported by hash functions.
- Operations on unordered containers (e.g., find, insert, delete) are $O(1)$.



std::unordered_[multi] set/map and Hash functions



std::unordered_[multi] set/map and Hash functions

s.insert("Pluto");

"first hash"
 $h = \text{std::hash<std::string>}$

$h(\text{"Pluto"})$

3080197140078183764

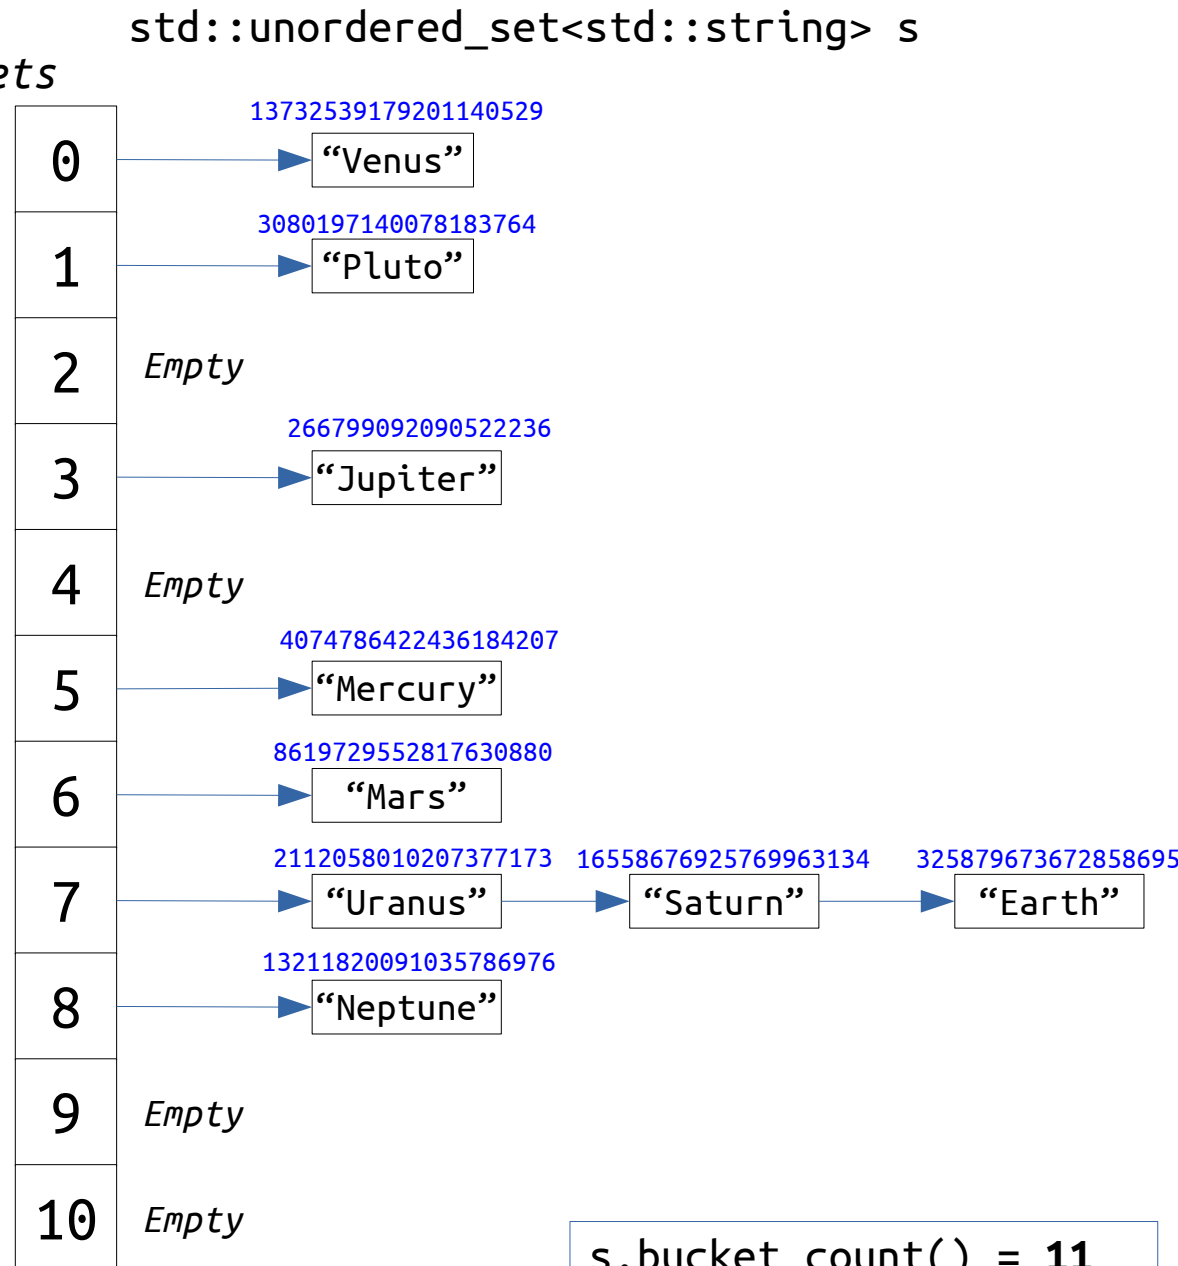
"second hash"
In which bucket insert?

$3080197140078183764 \% \text{s.bucket_count}()$

1

Insert "Pluto"
in bucket 1

buckets




s.bucket_count() = 11
s.size() = 9
s.load_factor() = 0.81

std::unordered_[multi] set/map and Hash functions

std::unordered_set<std::string> s

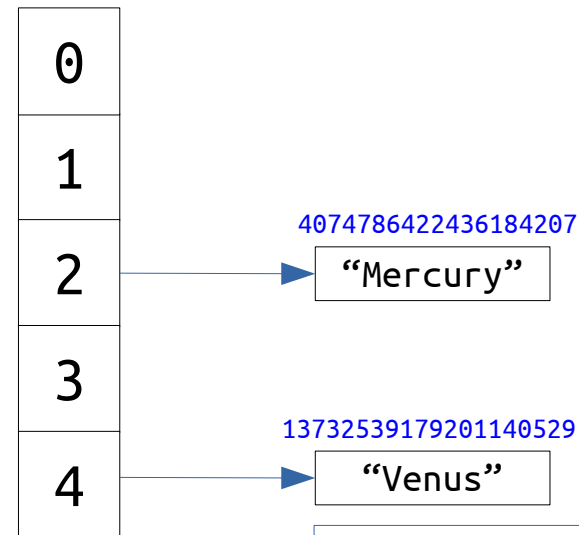
buckets



s.bucket_count() = 2
s.size() = 1
s.load_factor() = 0.5
s.max_load_factor() = 1



buckets



s.bucket_count() = 5
s.size() = 2
s.load_factor() = 0.4
s.max_load_factor() = 1

- When inserting a new element, if `load_factor() > max_load_factor()` then the table is enlarged (bucket_count)
- The *second* hash function "to know in which bucket insert" changes, so elements are re-allocated!

4074786422436184207 % s.bucket_count()

The *first* hash
`std::hash<std::string>`
does not change!